

The Reachability-Bound Problem

Sumit Gulwani

Microsoft Research
sumitg@microsoft.com

Florian Zuleger*

TU Vienna
zuleger@forstye.tuwien.ac.at

Abstract

We define the *reachability-bound problem* to be the problem of finding a symbolic worst-case bound on the number of times a given control location inside a procedure is visited in terms of the inputs to that procedure. This has applications in bounding resources consumed by a program such as time, memory, network-traffic, power, as well as estimating quantitative properties (as opposed to boolean properties) of data in programs, such as information leakage or uncertainty propagation.

Our approach to solving the reachability-bound problem brings together two different techniques for reasoning about loops in an effective manner. One of these techniques is an abstract-interpretation based iterative technique for computing precise disjunctive invariants (to summarize nested loops). The other technique is a non-iterative proof-rules based technique (for loop bound computation) that takes over the role of doing inductive reasoning, while deriving its power from the use of SMT solvers to reason about abstract loop-free fragments.

Our solution to the reachability-bound problem allows us to compute precise symbolic complexity bounds for several loops in .Net base-class libraries for which earlier techniques fail. We also illustrate the precision of our algorithm for disjunctive invariant computation (which has a more general applicability beyond the reachability-bound problem) on a set of benchmark examples.

Categories and Subject Descriptors C.4 [Performance of Systems]: Measurement techniques; Reliability, availability, and serviceability; D.2.4 [Software Engineering]: Software/Program Verification; F.3.1 [Logics and Meanings of Programs]: Specifying and Verifying and Reasoning about Programs

General Terms Verification, Performance, Reliability

Keywords Resource Bound Analysis, Disjunctive Invariants, Transitive Closure, Ranking Functions, Pattern Matching

1. Introduction

Program execution makes use of physical resources, and it is often important to compute worst-case bounds on usage of those resources as a function of the program inputs. For example, in

* The research of the second author was performed during an internship at MSR Redmond and was supported in part by Microsoft Research through its PhD Scholarship Programme.

memory-constrained environments such as embedded systems, it is important to bound the amount of *memory* required to run certain applications. In real-time systems, it is important to bound the worst-case *execution-time* of the program. Applications running on low-power devices or low-bandwidth environments must use up little *power* or *bandwidth* respectively. With the advent of cloud computing, where users would be charged per program execution, predicting resource usage characteristics would be a crucial component of *accurate bid placement* by cloud providers. One of the fundamental questions that needs to be answered for computing such resource bounds is: How many times is a given control location inside the program that consumes these resources executed?

Program execution also affects certain quantitative properties of data that it operates on. For example, how much *secret* information is leaked by a program depends on the number of times a certain operation that leaks the data, either by direct or indirect information flow, is executed [23]. Or the amount of *perturbation* in the output data values resulting from a small perturbation or uncertainty in the input values depends on the number of times additive error propagation operators are applied. This is the quantitative version of the boolean problem of continuity studied in [7]. Estimating such quantitative properties again requires addressing a similar question as above: How many times is a given control location inside the program that performs certain operations executed?

We refer to the problem of bounding the number of visits to a given control location π as the *reachability-bound problem*. Our two-step solution to this problem brings together two different techniques for reasoning about loops: an iterative technique for computing disjunctive invariants, and a non-iterative proof-rule based technique for computing bounds of transition systems.

The first step consists of generating a disjunctive *transition-system* that describes relationships between values of program variables that are live at π and their values in the immediate next visit to π . This requires summarizing inner loops that lie on a path from π back to itself for which we present an abstract interpretation based iterative algorithm that generates disjunctive loop invariants. The precision of our algorithm relies on a *convexity-like assumption*, which appears satisfied by all instances that we came across in practice, and leads to an interesting completeness theorem (Theorem 12). We also evaluated the precision of this algorithm on benchmark examples taken from recent work on computing disjunctive invariants. Our algorithm can discover required invariants in all examples, suggesting its potential for effective use in other applications requiring disjunctive invariants besides bound analysis.

The second step consists of generating bounds for the disjunctive transition system thus generated. For this, we propose a non-iterative proof-rules based technique that requires discharging queries using an off-the-shelf SMT solver. These proof rules describe conditions that are sufficient for combining the ranking functions for individual transitions into an overall bound of the transition system using three different mathematical operators, namely max, sum, and product. This is unlike existing work [4, 8, 28]

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

PLDI'10, June 5–10, 2010, Toronto, Ontario, Canada.

Copyright © 2010 ACM 978-1-4503-0019-3/10/06...\$10.00

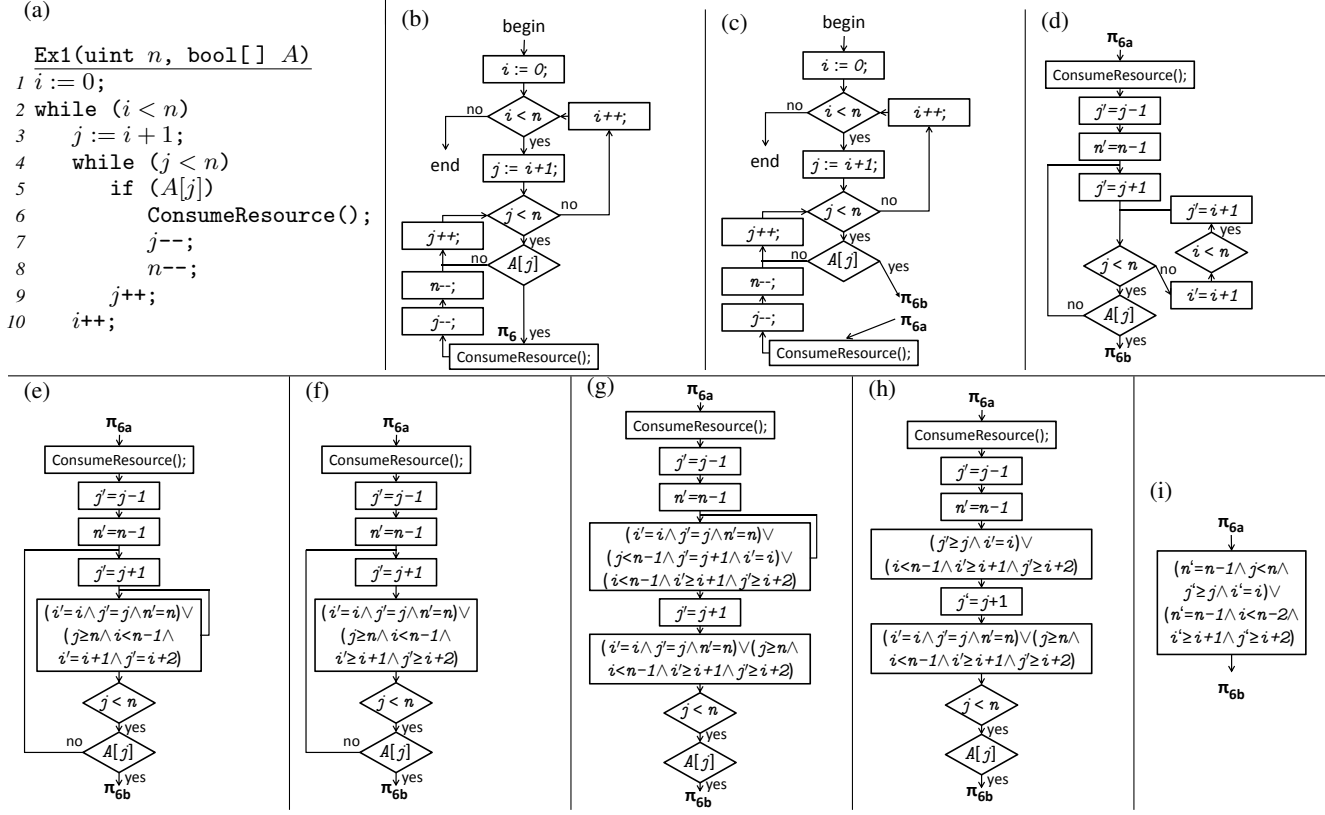


Figure 1. This figure illustrates generation of transition system for a given control location. (a) A loop skeleton from .Net base-class library. (b) Flow-graph representation of the program in (a). (c) Flow-graph obtained from (b) by splitting location π_6 into π_{6a} and π_{6b} . (d) Part of the flow-graph from (c) between π_{6a} and π_{6b} after re-drawing it. (e) Body of inner loop in (d) replaced by its transition system representation T_1 . (f) Inner loop in (e) replaced by transitive closure T'_1 of transition system T_1 . (g) Body of outer loop in (f) replaced by its transition system representation T_2 . (h) Outer loop in (g) replaced by Transitive closure T'_2 of transition system T_2 . (i) Transition system for π_6 .

on termination analysis where the goal is to generate any ranking function for a transition system with disregard to the precision of the ranking function. This methodology represents an interesting design choice for reasoning about loops, because SMT solvers are used to perform precise reasoning about transitions (loop-free code-fragments), whereas a simple proof-rules based technique takes over the role of performing inductive reasoning effectively. It will be interesting to consider applying such a methodology to other problems.

We have implemented our solution to the reachability-bound problem in a tool that computes symbolic computational complexity bounds for procedures in .Net codebases. This involves computing amortized complexity for nested loops by solving the reachability-bound problem for nested loops. To our knowledge, our analysis is the first that addresses the problem of computing the amortized complexity for nested loops. Existing techniques for bound analysis [16, 19, 15, 2] do not compute amortized complexity of nested loops, but instead over-approximate it by the product of the iterations of the outer loop and the worst-case complexity of the inner loop for any iteration of the outer loop, thereby leading to imprecise bounds.

Contributions and Organization

- We define the reachability-bound problem and the notion of a precise solution to that problem (Section 3). This contributes to the problem of defining an entire quantitative logic, which is part of the *quantitative agenda* set forth recently [21] (as opposed to the Boolean agenda).

- We describe an algorithm for the generation of a transition system based on transformations on reducible flowgraphs for reducing the problem of computing the reachability-bound to the problem of computing the bound for a transition system (Section 4).
- We describe an abstract interpretation based iterative algorithm for computing the transitive closure of a transition system, or, equivalently, disjunctive invariants for a loop. (Section 5).
- We describe non-iterative proof rules (Section 7) that allow computing precise symbolic bounds for a transition system from the ranking functions of individual transitions, which can be obtained using the technique described in Section 6.
- We present experimental results evaluating the effectiveness of various aspects of our solution (Section 8).

2. Motivating Examples and Technical Overview

In this section, we discuss some examples that are representative of some challenges that arise during the computation of symbolic bounds for the reachability-bound problem. We also provide a technical overview of our solution.

2.1 Bounding number of visits to a given control location

Consider the program shown in Figure 1, and consider the problem of computing symbolic bounds on the number of times the procedure $\text{ConsumeResource}()$ is called at Line 6. One approach would be to approximate it by computing a bound on the number of iterations of the closest enclosing loop at Line 4 using techniques for

loop bound computation (as in [16, 19]). However, this approach will yield quite conservative results since the number of iterations of the loop at Line 4 is bounded above by n^2 , while the number of executions of Line 6 is bounded above by n .

Our approach first computes the following symbolic relationship between values of variables i, j, n at Line 6 with their values i', j', n' in the immediate next visit to Line 6. The relationship is expressed as the disjunction of two transitions given as two disjuncts:

$$(n' = n - 1 \wedge j < n - 1 \wedge j' \geq j \wedge i' = i) \\ \vee (n' = n - 1 \wedge i < n - 1 \wedge i' \geq i + 1 \wedge j' \geq i + 2)$$

This is done using the `GenerateTransitionSystem` algorithm described in Figure 4 in Section 4. The algorithm enumerates all paths in the control-flow graph in Figure 1(d) between the locations π_{6a} and π_{6b} obtained after splitting the location π_6 in the original control-flow graph (in Figure 1(b)) into π_{6a} and π_{6b} (as shown in Figure 1(c)). (Note, that such a relationship is different from transition invariants [28] or variance assertions [4] that relate values of variables at a control location with their values in *any* successive iteration, as opposed to the immediate next iteration). The challenge in such an enumeration is that the number of paths in presence of loops between control locations π_{6a} and π_{6b} is not finite. For this purpose, the loops are summarized by disjunctive relationships between the inputs/outputs of the loop. These disjunctive relationships are generated by computing the transition system of the loop (recursively, using the same algorithm applied to the control location immediately after the loop-header) and then computing its transitive closure (using the algorithm `TransitiveClosure` described in Fig. 6 in Section 5). The transition system of the inner loop in Figure 1(d) is shown in Figure 1(e). The transitive closure of the inner loop is given in Figure 1(f). The transition system of the outer loop is shown in Figure 1(g). The transitive closure of the outer loop is given in Figure 1(h). The resulting transition system for location π_6 is given in Figure 1(i).

Next, bounds are computed for the transition system thus generated. This involves computing the ranking functions for the two transitions of the transition system for location (π_6) . A ranking function for a transition s is an integer-valued function (of the variables that occur in s) that is bounded below by 0 and decreases every time the transition s is taken. (For a formal definition see Section 6.) By using the pattern matching techniques described in Section 6 we compute the ranking functions $n - j$ resp. $n - i - 2$ for the transition given by the first resp. second disjunct of the transition system for (π_6) . These ranking functions are then composed using one of the proof rules described in Section 7 (in this case, the proof rule in Theorem 16) to obtain a bound of $\text{Max}(0, n - j, n - i - 2)$ in terms of the inputs to the transition system (For details, see Example 17). Using the invariants $i \geq 0 \wedge j \geq 1$ that hold during the first visit to π_6 (which can be obtained by generating invariants at control location π_{6b} in Figure 1(d)), we obtain a bound of $n - 1$ on the transition system in terms of procedure inputs. This implies a bound of n on the number of visits to control location π_6 .

2.2 Bounding iterations of a loop

Computing bounds on the number of loop iterations is a special case of the reachability-bound problem where the control location under consideration is the location immediately after the loop header. Under that case, our technique outperforms recent techniques for loop bound computation and termination.

In particular, our technique is able to compute the bounds for loops whose iterations are affected by inner loops for which existing bound techniques (such as [2, 16, 19]) mostly fail (for details, see related work in Section 9). Such loops are quite common in .Net base-class library, and Figure 2 gives some examples. One of

Ex6(int n, int x, int z)	Ex7(uint n, uint m)
1 while (x < n)	1 Assume(0 < n < m);
2 if (z > x) x++;	2 j := n + 1;
3 else z++;	3 while (j < n ∨ j > n)
	4 if (j > m) j := 0;
	5 else j++;

Figure 3. Loop templates Ex6 and Ex7 (from Microsoft product-code) taken respectively from recent work on proving termination [8] and loop bound computation [16]. Our proof rules for bound computation provide an alternative, but simpler, formalism for computing bounds. (For details see Example 20 and Example 23.)

the key challenges addressed by our technique in such examples is the summarization of the inner loops by precise transitive-closure of the transition system represented by these loops (in effect, disjunctive relationships between the inputs and outputs of the loop).

Also, even in case of loops with no nested loops, our technique is able to compute bounds for loops using a much simpler uniform algorithm compared to existing termination techniques or specialized bound computation techniques. Figure 3 shows two such examples that have been used as motivating examples by previous techniques. The computation of the transition system for these examples is almost trivial, and the bound computation of the resulting transition system is enabled by simple but precise proof rules (Theorem 19 and Theorem 21) for bound computation from ranking functions of individual transitions (for details, see Example 20 and 23).

3. The Reachability-Bound Problem

There are two classical problems associated with the reachability of a control location π inside a procedure P with inputs \vec{n} .

- **Safety Problem:** Is the control location π never reached/visited?
- **Termination Problem:** Is the control location π visited at most a finite number of times?

In this paper, we have motivated the following bound problem, which is a more general instance of both the safety and termination problem. In fact, as we will see, our solution to the bound problem builds over techniques for safety and termination checking.

- **Reachability-bound Problem:** Compute a worst-case symbolic bound $\mathcal{B}(\vec{n})$ on the number of visits to π for any execution of P .

The notion of a worst-case symbolic bound is defined below.

DEFINITION 1 (Worst-case symbolic bound). An integer-valued function $\mathcal{B}(\vec{n})$ is a worst-case symbolic bound for a control location π inside a procedure P with inputs \vec{n} if for any input state \vec{n}_0 , the number of times π is visited is at most $\mathcal{B}(\vec{n}_0)$.

There may be multiple worst-case symbolic bounds for a given location. It is desirable to produce a bound that is *precise* in the sense that there exists a family $\phi(\vec{n})$ of worst-case inputs that exhibit the bound (up to some constant factor, as motivated by the definition of asymptotic complexity), formally defined as follows:

DEFINITION 2 (Precision of a worst-case symbolic bound). A worst-case symbolic bound $\mathcal{B}(\vec{n})$ for a control location π inside a procedure P with inputs \vec{n} is said to be precise (up to multiplicative constant factors) if there exist positive integers c_1, c_2 , and a formula $\phi(\vec{n})$ such that:

- E1. For any assignment \vec{n}_0 to variables \vec{n} such that $\phi(\vec{n}_0)$ holds, the number of times control location π is visited (when procedure P is executed in the input state \vec{n}_0) is at least $\frac{\mathcal{B}(\vec{n}_0)}{c_1} - c_2$.

Ex2(uint n , uint m)	Ex3(uint n , bool[] A)	Ex4(uint n)	Ex5(uint n)
1 while ($n > 0 \wedge m > 0$) 2 $n--$; $m--$; 3 while (nondet()) 4 $n--$; $m++$;	1 while ($n > 0$) 2 $t := A[n]$; 3 while ($n > 0 \wedge t = A[n]$) 4 $n--$;	1 flag := true; 2 while (flag) 3 flag := false; 4 while ($n > 0 \wedge$ nondet()) 5 $n--$; flag := true;	1 $i := 0$; 2 while ($i < n$) 3 flag := false; 4 while (nondet()) 5 if (nondet()) 6 flag:=true; $n--$; 7 if (\neg flag) $i++$;
$n' \leq n \wedge m' \geq m$	$n' \leq n$	$(n' \leq n-1 \wedge \text{flag}')$ $\vee (\text{Same}(\{n, \text{flag}\}))$	$(n' \leq n-1 \wedge \text{flag}' \wedge i' = i)$ $\vee (\text{Same}(\{i, n, \text{flag}\}))$
$n > 0 \wedge m > 0 \wedge n' \leq n-1$	$(n > 0 \wedge n' \leq n \wedge A[n] \neq A[n'])$ $\vee (n > 0 \wedge n' \leq 0)$	$(\text{flag} \wedge \text{flag}' \wedge n > 0 \wedge n' \leq n-1)$ $\vee (\text{flag} \wedge \neg \text{flag}' \wedge n' = n)$	$(i < n \wedge \text{flag}' \wedge n' \leq n-1 \wedge i' = i)$ $\vee (i < n \wedge \neg \text{flag}' \wedge i' \geq i+1 \wedge n' = n)$
n	n	$n+1$	n

Figure 2. Loop templates from .Net class libraries where iterators of a loop are modified by inner loops. The second row shows the required transitive closure of the inner loops to enable precise symbolic bound computation of respective outer loops. The third row shows the resultant transition-system generated for the outer loops after summarizing the respective inner loops by the transitive closure of their transition-system (using the algorithm in Figure 4). The final (fourth) row shows the bound computed from the transition-system by the algorithm in Figure 7. We use the predicate $\text{Same}(V)$ inside a transition to denote that the variables in V do not change their value, i.e., $\text{Same}(V) = \bigwedge_{x \in V} (x' = x)$.

E2. For any integer k , there exists a satisfying assignment \vec{n}_1 for $\phi(\vec{n})$ such that $\mathcal{B}(\vec{n}_1) > k$. In other words, the formula $\exists \vec{n} : (\mathcal{B}(\vec{n}) \geq k \wedge \phi(\vec{n}))$ has a satisfying assignment.

We refer to the triple (ϕ, c_1, c_2) as precision-witness for bound \mathcal{B} .

The following example explains and motivates the requirements E1 and E2 in the above definition.

EXAMPLE 3. A precision-witness for the bound of n on the number of times Line 6 is visited in the example program Ex1 in Figure 1 can be $\phi = \forall k(0 \leq k < n \Rightarrow A[k])$, $c_1 = 1$ and $c_2 = 1$ since it can be shown that under the precondition ϕ , Line 6 is visited at least $n-1$ times.

A precision-witness for the bound of n^2 on the number of times the inner loop (Line 5) is executed can be $\phi = \forall k(0 \leq k < n \Rightarrow \neg A[k])$, $c_1 = 4$ and $c_2 = 1$ since it can be shown that under the precondition ϕ , Line 5 is visited at least $n^2/4$ times. This is because, for example, i takes all values between 0 to $n/2-1$ at Line 2 (hence the number of visits to Line 2 is at least $n/2$), and for each of those visits, j takes all values between $n/2$ to $n-1$ at Line 4 (i.e., the number of visits to Line 4 is at least $n/2$). Note that if we did not relax the requirement E1 to allow for constants c_1 and c_2 , then computation of a precise bound would have required us to compute the exact bound of $\frac{(n-1)(n-2)}{2}$. It would be impractical to find such exact closed-form solutions.

A bound of, say, 100, on the number of times Line 6 is visited is not precise. It may appear that $\phi = (\forall k(0 \leq k < 100 \Rightarrow A[k])) \wedge n \leq 100$, $c_1 = 1$ and $c_2 = 1$ is a precision-witness. However, note that it violates requirement E2 since for $k = 101$ (in fact, for any k greater than 100), there does not exist a satisfying assignment for the formula $\phi \wedge 100 \geq 101$.

In this paper, we describe an algorithm for computing a worst-case symbolic bound. Manual investigation of the bounds returned by our algorithm on our benchmark examples confirms that the bounds are precise. Automatically establishing the precision of a bound \mathcal{B} returned by our algorithm is an orthogonal problem that we are currently working on. It requires identifying a precision-witness (ϕ, c_1, c_2) and establishing that $\frac{\mathcal{B}}{c_1} - c_2$ is a lower bound for all inputs that satisfy ϕ . The duality between the problems of computing a symbolic bound \mathcal{B} and the problem of finding a witness ϕ to show that \mathcal{B} is precise is similar to the duality between the problems of proving a given safety property, or finding a concrete counterexample/witness to the violation of a safety property. However, the challenge in our case is that the witness ϕ that establishes the precision of a given symbolic bound is symbolic as opposed to being concrete.

We next describe our overall algorithm for bound computation.

3.1 Algorithm

Our algorithm for the reachability-bound problem is as follows.

```

ReachabilityBound( $\pi$ )
1   $T := \text{GenerateTransitionSystem}(\pi)$ ;
2   $\mathcal{B} := 1 + \text{ComputeBound}(T)$ ;
3  return  $\text{TranslateBound}(\mathcal{B}, \pi)$ ;

```

Line 1 of the algorithm first computes a disjunctive transition system T for the control location π that describes how the variables at π get updated in the immediate next visit to control location π . This is done using the algorithm described in Figure 4 (Section 4), which in turn uses the algorithm for transitive closure computation described in Figure 6 (Section 5) to summarize any inner loops.

Line 2 of the algorithm computes a bound for the transition system T using the algorithm described in Figure 7 (Section 7), which in turn makes use of techniques described in Section 6 for computing ranking functions of individual transitions. The bound \mathcal{B} on number of visits to π is then obtained by adding 1 to the bound for transition system T to account for the first visit to π .

The bound \mathcal{B} is expressed in terms of inputs to the transition system, which may not necessarily be the procedure inputs. The function TranslateBound at Line 3 then translates the bound \mathcal{B} at π in terms of the procedure inputs. This can be done either by using invariants (computed with an invariant generation tool) that relate the procedure inputs with the inputs to the transition system T , or by using a backward symbolic engine to express the transition system inputs in terms of the procedure inputs. We implemented the latter approach, which we found to be extremely effective in terms of both precision and efficiency. This technique is detailed in [17].

Notice, how our solution builds on techniques for safety or termination checking. Step 1 uses disjunctive invariants, which is essentially what is needed for safety checking. Step 2 uses ranking functions, which are required for termination checking. Use of these techniques together with novel proof-rules for composing ranking functions yields an effective solution to the bound problem.

4. Generation of Transition System

We first define the notion of a transition and a transition system with regard to a control location π .

DEFINITION 4 (Transition for a Control Location π). Let \vec{x} be the tuple of the variables live at π . A transition for π is a relation $T(\vec{x}, \vec{x}')$ between variables \vec{x} and their primed counterparts \vec{x}' such that if \vec{x} take values \vec{v}_1 and \vec{v}_2 during any two immediate successive/consecutive visits to π , then $T(\vec{v}_1, \vec{v}_2)$ holds.

A transition is always assumed to be represented as a conjunction of formulae over the variables \vec{x} and \vec{x}' .

DEFINITION 5 (Transition System for a Control Location π). A transition system is a set of transitions.

A transition system is always assumed to be represented as a DNF formula where every disjunct corresponds to the representation of a transition of the transition system.

We desire a disjunctive representation for our transition system since our bound computation algorithm in Section 7 works by identifying precise ranking functions for a single transition/path, and then using proof rules to obtain the ranking function/bound for the entire transition system.

The key idea for generating a transition system for a control location π is to split the control location π into the two locations (π_a, π_b) (using the `Split` transformation shown in Figure 5(a)) and enumerate all paths that start at π_a and end at π_b and take the disjunctions of the transitions represented by each path. The challenge that arises in such an enumeration is the presence of any nested loops. We address this challenge by replacing the nested loop by the transitive closure of the transition system of the nested loop (using the `Summarize` transformation shown in Figure 5(b)). Since path enumeration leads to an exponential blowup, we generate the transition systems on the flowgraph that has been sliced with respect to the statements on which π is control-dependent [25] (since these are the statements that determine the number of times π is executed). This usually leads to transition systems with a very small number of transitions, as is exemplified by statistics in Fig. 8 (Section 8.1).

Figure 4 describes the algorithm to generate the transition system for a control location π . The algorithm is described at flowgraph level. We make the assumption about the flowgraphs being reducible, but not necessarily structured. Our algorithm can be extended to irreducible flowgraphs too; but we avoid that for ease of presentation, and the fact that most flowgraphs in practice are in fact reducible [25]. However, it is important to consider the case of unstructured flowgraphs because even if the original flowgraph was structured, after the splitting transformation, the new flowgraph would no longer be structured. The splitting transformation, however, is reducibility-preserving.¹

Line 1 transforms the flowgraph by splitting the input control location π into two locations π_a and π_b using the `Split` transformation described in Figure 5(a). The loop in Line 2 iterates over each top-level loop L in the transformed flowgraph. (Recall that any graph can be decomposed into a DAG of maximal strongly-connected components.) Line 3 makes use of the fact that every loop in a reducible flow-graph has a unique header node. Line 4 recursively generates the transition system for the loop L in the transformed flow-graph, while Line 5 generates its transitive closure (using the algorithm described in Figure 6 in Section 5). Lines 6 and 7 replace the loop L by its summary obtained by generating transitive closure of the transition system represented by it (using the `Summarize` transformation shown in Figure 5(b)). The effect of the foreach-loop in Line 2 is to replace all loops on the paths between π_a and π_b by (disjunctive) loop-free abstract code-fragments. The transition system can now simply be generated by enumerating all paths (which are now finite in number) between π_a and π_b .

Lines 8-10 generate the transition system for an acyclic flowgraph by a simple forward dataflow analysis that associates a (disjunctive) transition system $F[\pi]$ with each edge/control location π in the transformed flowgraph. For this purpose, we associate the

```

GenerateTransitionSystem( $\pi$ )
1 ( $\pi_a, \pi_b$ ) := Split( $\pi$ );
2 foreach top-level loop  $L$ :
3    $\pi_L$  := location before header of  $L$ ;
4    $T$  := GenerateTransitionSystem( $\pi_L$ );
5    $T_c$  := TransitiveClosure( $T$ );
6   Insert Summary( $T_c$ ) before header;
7   Remove back-edges;
8 Initialize  $F[\pi_a]$  to the transition system Id;
9 Propagate transitions  $F$  using Merge/Compose rules;
10 return  $F[\pi_b]$ ;

```

Figure 4. Generation of transition system for a control location π .

entry location π_a with the transition system consisting of a single transition `Id`, which is the identity mapping between the variables and their primed versions. The transfer functions for performing this dataflow analysis are described in Figure 5. Without loss of any generality, we assume that all conditional guards have been translated into `Assume` statements. The `Merge` transfer function simply returns the disjunctions of the transitions in the two input transition systems. The `Compose` transfer function makes use of the compose operator \circ that returns the composition of two transitions.

DEFINITION 6 (Composition of Transition Systems). Given two transition systems $T(\vec{x}, \vec{x}') = \bigvee_i s_i$ and $T'(\vec{x}, \vec{x}') = \bigvee_j s'_j$, we define their binary composition to be

$$T \circ T' \stackrel{def}{=} \bigvee_{i,j} s_i \circ s'_j,$$

where $s_i \circ s'_j$ denotes the transition

$$s_i(\vec{x}, \vec{x}') \circ s'_j(\vec{x}, \vec{x}') \stackrel{def}{=} \exists \vec{x}'' (s_i[\vec{x}''/\vec{x}'] \wedge s'_j[\vec{x}''/\vec{x}']),$$

where $s_i[\vec{x}''/\vec{x}']$ denotes the substitution of \vec{x}' by \vec{x}'' in s_i .

The `Translate` function converts a statement into a transition system as follows. Without loss of any generality, we assume that the only assignment statement is of the form $x := e$ since memory can be modeled using `Select` and `Update` expressions. The other kinds of statements can be either an `Assume` statement (obtained from the conditional guards) or a `Summary` statement (obtained from the summarization of nested loops).

$$\text{Translate}(x := e) = (x' = e) \wedge \left(\bigwedge_{y \neq x} y' = y \right)$$

$$\text{Translate}(\text{Assume}(\text{guard})) = \text{Id} \wedge \text{guard}$$

$$\text{Translate}(\text{Summary}(T)) = T$$

EXAMPLE 7. The transition system for control location π_6 in Figure 1(b) is shown in Figure 1(e) along with the various steps required to obtain it from the flowgraph in Figure 1(d). These include computing the transition system for the inner loop and then replacing the inner loop by its transitive closure. Next, the process is repeated for the outer loop.

5. Computation of Transitive Closure

In this section, we describe an algorithm for computing a transitive closure (defined below) of a transition system. This operation is required by the `GenerateTransitionSystem` algorithm described in Figure 6 in the previous section.

DEFINITION 8 (Transitive Closure). We say that $T'(\vec{x}, \vec{x}')$ is a transitive closure of a transition system $T(\vec{x}, \vec{x}')$ if

$$\text{Id} \Rightarrow T' \quad \text{and} \quad T' \circ T \Rightarrow T'$$

¹ It is interesting to observe that the nesting structure of the loops inside which π was originally nested, is completely reversed after the splitting transformation, but the flowgraph stays reducible.




Figure 5. This figure describes the flowgraph transformations *Split* and *Summarize*, and the transfer functions *Compose* and *Merge* required in the algorithm *GenerateTransitionSystem* for computing the transition system for any control location.

EXAMPLE 9. Figure 1(e) provides an example of a transition system T and its transitive closure. Note that $i' \geq i$ is another choice for the transitive closure for T . However, it is not as precise as the one shown in Figure 1(e), and would lead to the generation of a transition system for location π_6 for which no bound exists.

Generating the transitive closure of a transition system is like computing the invariants for a loop which represents the transition system. Example 9 suggests the importance of these invariants to be precise, and hence disjunctive. There has been some work on discovering disjunctive invariants [5, 16, 20, 29, 13, 14] in general. We present below a technique that takes advantage of its particular application to bound analysis. (We also remark that our technique can be used in general for proving safety properties of programs. In Section 8.2, we present preliminary results that demonstrate the effectiveness of our technique on a set of benchmark examples taken from a variety of recent literature on generating disjunctive invariants.)

Our algorithm for the computation of precise transitive closures is inspired by a *convexity-like* assumption that we found to hold true for all examples we have come across in practice. (This includes the desired transitive closure of the transitions-systems of nested loops to compute precise bounds, as well as the benchmarks considered by previous work on computing disjunctive invariants.)

Recall that a theory is said to be *convex* iff for every quantifier-free formula ϕ in that theory, if ϕ implies a disjunction of equalities, then it implies one of those equalities, i.e.,

$$\left(\phi \Rightarrow \left(\bigvee_i (x_i = y_i) \right) \right) \Rightarrow \left(\bigvee_i (\phi \Rightarrow (x_i = y_i)) \right) \quad (1)$$

Now, if $\bigvee_{j=1}^m s'_j$ is a transitive closure of $\bigvee_{i=1}^n s_i$, then it follows from the definition of the transitive closure, that for all $i \in \{1, \dots, n\}$ and $j \in \{1, \dots, m\}$, the following holds:

$$Id \Rightarrow \bigvee_{k=1}^m s'_k \quad \text{and} \quad s'_j \circ s_i \Rightarrow \bigvee_{k=1}^m s'_k$$

After distributing implication over disjunctions in the above equations (in a manner similar to in Equation 1), we obtain the *convexity-like* assumption, which is defined formally below.

DEFINITION 10 (Convexity-like Assumption).

Let $T' = \bigvee_{j=1}^m s'_j(\vec{x}, \vec{x}')$ be a transitive closure for a transition system $T = \bigvee_{i=1}^n s_i(\vec{x}, \vec{x}')$, where each s_i and s'_j is a conjunctive relation. We say that the transitive closure $\bigvee_j s'_j$ satisfies the convexity-like assumption if there exists an integer $\delta \in \{1, \dots, m\}$, a map $\sigma : \{1, \dots, m\} \times \{1, \dots, n\} \mapsto \{1, \dots, m\}$, such that for all $i \in \{1, \dots, n\}$ and $j \in \{1, \dots, m\}$, the following holds:

$$Id \Rightarrow s'_\delta \quad \text{and} \quad (s'_j \circ s_i) \Rightarrow s'_{\sigma(j,i)}$$

```

TransitiveClosure( $\bigvee_{i=1}^n s_i$ )
1  for  $j \in \{1, \dots, m\} - \{\delta\}$ :  $s'_j := \text{false}$ ;
2   $s'_\delta := Id$ ;
3  do {
4    for  $i \in \{1, \dots, n\}$  and  $j \in \{1, \dots, m\}$ :
5       $s'_{\sigma(j,i)} := \text{Join}(s'_{\sigma(j,i)}, s'_j \circ s_i)$ 
6  } while any change in  $\bigvee_{j=1}^m s'_j$ 
7  return  $\bigvee_{j=1}^m s'_j$ ;

```

Figure 6. Transitive closure computation of a transition system.

The tuple (δ, σ) is referred to as a convexity-witness of $\bigvee_{j=1}^m s'_j$.

The convexity-like assumption essentially implies that no case-split reasoning is needed to prove inductiveness of transitive closure.

EXAMPLE 11. All the transitive closures of the respective transition systems described in Figure 1(e) and Figure 2 satisfy the convexity-like assumption. For example, the convexity-witness for the transitive closure of the transition system T shown in Figure 1(e) is $\delta = 1$ and $\sigma = \{(1, 1) \mapsto 2, (2, 1) \mapsto 2\}$. A convexity-witness for the transitive closure of the transition system T' shown in Figure 1(e) is $\delta = 1$ and $\sigma = \{(1, 1) \mapsto 1, (2, 1) \mapsto 2, (1, 2) \mapsto 2, (2, 2) \mapsto 2\}$.

Given a convexity-witness (δ, σ) of any transitive-closure T' (that satisfies the convexity-like assumption) of a transition system T , the algorithm in Figure 6 describes a way to compute a transitive closure that is at least as precise as T' . This property (stated formally in the following theorem) is quite significant in light of the fact that discovering disjunctive invariants has been quite a challenging task in literature and several merging heuristics based on semantics of the constituent dataflow facts have been suggested. The following theorem states the remarkable result that a semantic merging criterion cannot be better than a static syntactic criterion for merging data-flow facts.

THEOREM 12 (Precision of TransitiveClosure Algorithm).

Let $\bigvee_{j=1}^m s''_j$ be any transitive closure of a given transition system $\bigvee_{i=1}^n s_i$ that satisfies the convexity-like assumption. Given the number of disjuncts m and a convexity-witness (δ, σ) , algorithm in Figure 6 outputs a transitive closure that is at least as precise as $\bigvee_{j=1}^m s''_j$.

PROOF: We can prove that $s'_j \Rightarrow s''_j$ by induction on the number of loop iterations; the base case as well as the inductive case both follows easily from the definition of *convexity-like* assumption. \square

The algorithm in Figure 6 performs abstract interpretation over the power-set extension of an underlying abstract domain (such as polyhedra [10], octagons [24], conjunctions of a given set of predicates), where the elements are restricted to at most m disjuncts. We assume, that the underlying abstract domain is equipped with a Join operator, which takes two elements and returns the least upper bound of both elements. The algorithm uses the map σ to determine how to merge the $n \times m$ different disjuncts (into m disjuncts) that are obtained after the propagation of m disjuncts across n transitions using the Join operator. The key distinguishing feature of the algorithm from earlier work on computing disjunctive invariants is that our algorithm uses a syntactic criterion based on σ to merge disjuncts as opposed to using a semantic criterion based on the notion of differences between disjuncts. This is justified by Theorem 12, which, in effect, says that no semantic merging criterion can be more powerful than a static syntactic criterion. There are two issues with the algorithm presented in Figure 6 that we discuss below.

Abstract Domains with Infinite Height The algorithm may not on domains with infinite height. The standard solution would be to the apply a Widen operator (as defined in [9]) in place of the Join operator, in order to enforce termination.

Since the use of widening may overapproximate the least fixed point in general, it is no longer possible to formally prove precision results as in Theorem 12. However, we show experimentally (in Section 8.2) that our algorithm is able to compute precise enough invariants with the use of standard widening techniques when applied on benchmarks taken from recent work on computing disjunctive invariants.

Choice of m and a convexity-witness (δ, σ) Since we do not know the desired transitive closure and its convexity-witness (δ, σ) upfront, we have two options.

Option 1: We can enumerate all possible (δ, σ) for a specifically chosen m . There are m^{mn} such possible maps since without loss of any generality, we can assume that δ is 1. If m and n are small constants, say 2 (which is quite often an important special case), then there are 16 possibilities. Each choice for σ and δ results in some transitive closure computation by the algorithm. One can then select the strongest transitive closure among the various transitive closures thus obtained (or heuristically select between incomparable transitive closures). However, if m or n is large, then this approach quickly becomes prohibitive.

Option 2: We can use some heuristics to construct m, δ, σ . The following heuristic turns out to be the most effective for our application of bound computation. We set m and δ to $n + 1$, and select the map σ from the DAG of dependencies between transitions of T generated from bound computation of T (as described in Section 7). In particular, for any $i, j \in \{1, \dots, n\}$, we define $\sigma(n + 1, i) := i$, $\sigma(i, i) := i$, and $\sigma(i, j) := i$ except when $\neg \text{NI}(s_j, s_i, r)$ (where $r \in \text{RankC}(s_i)$ was the ranking function that contributed to the bound computation of T) in which case we define $\sigma(i, j) := j$. It can be proved that such a choice of the map δ and σ would generate a transitive closure that would allow for computing the bound of $(T \circ \text{TransitiveClosure}(T))$ using the bound computation algorithm described in Section 7, provided it was able to generate a bound for the transition system T . Such a transitive closure preserves important relationships (between the program variables) for the application of computing the bound of the transition system that is to be obtained after replacing the corresponding loop by the transitive closure. In particular, note that this heuristic for the construction of a convexity-witness, when used in conjunction with the algorithm in Figure 6 discovers the required transitive-closures of the respective transition systems mentioned in Figure 1(e) and Figure 2.

6. Ranking Function for a Transition

In this section, we show how to compute a *ranking function* for a transition. These ranking functions are made use of by the bound computation algorithm described in Section 7.

DEFINITION 13 (Ranking Function for a Transition). We say that an integer-valued function $r(\vec{x})$ is a ranking function for a transition $s(\vec{x}, \vec{x}')$ if it is bounded below by 0 and if it decreases by at least 1 in each execution of the transition, i.e.,

- $s \Rightarrow (r > 0)$
- $s \Rightarrow (r[\vec{x}'/\vec{x}] \leq r - 1)$

We denote this by $\text{Rank}(s, r)$.

We say that a ranking function $r_1(\vec{x})$ is *more precise* than a ranking function $r_2(\vec{x})$ if $r_1 \leq r_2$ (because in that case, r_1 provides a more precise bound for the transition than r_2).

We discuss below the design of a functionality RankC that takes as input a transition $s(\vec{x}, \vec{x}')$ and outputs a set of ranking functions $r(\vec{x})$ for that transition. We use a pattern-matching based technique that relies on asking queries that can be discharged by an SMT solver. We found this technique to be effective (fast and precise) for most of the transitions that we encountered during the process of bound computation on .Net base-class libraries. However, other techniques, such as constraint-based techniques [27] or counter instrumentation enabled iterative fixed-point computation based techniques [15, 19] can also be used for generating ranking functions. Clearly, there are examples where the constraint-based or iterative techniques that perform precise arithmetic reasoning would be more precise, but nothing beats the versatility of simple pattern matching that can handle non-arithmetic patterns with equal ease.

We list below some patterns that we found to be most effective.

6.1 Arithmetic Iteration Patterns

One standard way to iterate over loops is to use an arithmetic counter. Ranking functions for such an iteration pattern can be computed using the following pattern.

If $s \Rightarrow (e > 0 \wedge e[\vec{x}'/\vec{x}] < e)$, then $e \in \text{RankC}(s)$

The candidates for expression e while applying the above pattern are restricted to those expressions that only involve variables from \vec{x} and those that occur syntactically as an operator of conditionals when normalized to the form $(e > 0)$, after rewriting a conditional of the form $(e_1 > e_2)$ to $(e_1 - e_2 > 0)$. In the following we give example transitions whose ranking functions can be computed using an application of this pattern.

- $\text{RankC}(i' = i + 1 \wedge i < n \wedge i < m \wedge n' = n \wedge m' \leq m) = \{n - i, m - i\}$
- $\text{RankC}(n > 0 \wedge n' \leq n \wedge A[n] \neq A[n']) = \{n\}$

The second example transition above (obtained from the transition system generated for the loop in the example program Ex3 in Figure 2) is a good illustration of how simple pattern matching is used to guess a ranking function, and an SMT solver (that can reason about combination of theory of linear arithmetic and theory of arrays) can be used to perform the relatively complicated reasoning of verifying the ranking function over a loop-free code fragment.

Another common arithmetic pattern is the use of a multiplicative counter whose value doubles or halves in each iteration (as in case of binary search). A more precise ranking function for such a transition can be computed by using the pattern below.

If $s \Rightarrow (e \geq 1 \wedge e[\vec{x}'/\vec{x}] \leq e/2)$, then $\log e \in \text{RankC}(s)$

The candidates for expression e while applying the above pattern are restricted to those expressions that only involve variables

from \vec{x} and those that occur syntactically as an operator of conditionals when normalized to the form $(e > 1)$, after rewriting a conditional of the form $(e_1 > e_2)$ that occurs in s to $(\frac{e_1}{e_2} > 1)$, provided e_2 is known to be positive. In the following we give example transitions whose ranking functions can be computed using an application of this pattern.

- $\text{RankC}(i' \leq i/2 \wedge i > 1) = \{\log i\}$
- $\text{RankC}(i' = 2 \times i \wedge i > 0 \wedge n > i \wedge n' = n) = \{\log(n/i)\}$

The above two patterns are good enough to compute ranking functions for most loops that iterate using arithmetic counters. However, for the purpose of completeness, we describe below two examples (taken from some recent work on proving termination) that cannot be matched using the above two patterns, and hence illustrate the limitations of pattern-matching. However, we can find ranking functions or bounds for these examples using the counter instrumentation and invariant generation techniques described in [15].

- Consider the terminating transition system $(x' = x + y \wedge y' = y + 1 \wedge x < n \wedge n' = n)$ from [6], which uses the principle of polyranking lexicographic functions for proving its termination. Note that the reason why the transition system terminates is because even though y is not known to be always positive, it will eventually become positive by virtue of the assignment $y' = y + 1$.
- Consider the terminating transition system $(x' = y \wedge y' = x - 1 \wedge x > 0)$. This transition system can be proven terminating by monotonicity constraints as introduced in [3]). Note, that the reason why the transition system terminates is because in every two iterations the value of x decreases by 1.

6.2 Boolean Iteration Patterns

Often loops contain a path/transition that is meant to execute just once. The purpose of such a transition is to switch between different phases of a loop, or to perform the cleanup action immediately before loop termination. Such an iteration pattern can be captured by the following rule/lemma, where the operator $\text{Bool2Int}(e)$ maps boolean values `true` and `false` to 1 and 0 respectively.

If $s \Rightarrow (e \wedge \neg(e[\vec{x}'/\vec{x}]))$, then $\text{Bool2Int}(e) \in \text{RankC}(s)$

The candidates for boolean expression e while applying the above pattern are restricted to those expressions that only involve variables from \vec{x} and those that occur syntactically in the transition s . In the following we give example transitions whose ranking functions can be computed using an application of this pattern.

- $\text{RankC}(\text{flag}' = \text{false} \wedge \text{flag}) = \{\text{Bool2Int}(\text{flag})\}$
- $\text{RankC}(x' = 100 \wedge x < 100) = \{\text{Bool2Int}(x < 100)\}$

6.3 Bit-vector Iteration Patterns

One standard way to iterate over a bit-vector is to change the position of the lsb, i.e., the least significant one bit (or msb, i.e., most significant one bit). Such an iteration pattern can be captured by the following rule/lemma, where the function $\text{LSB}(x)$ denotes the position of the least significant 1-bit, counting from 1, and starting from the most significant bit-position. $\text{LSB}(x)$ is defined to be 0 if there is no 1-bit in x . Note that $\text{LSB}(x)$ is bounded above by the total number of bits in bit-vector x .

If $s \Rightarrow (\text{LSB}(x') < \text{LSB}(x) \wedge x \neq 0)$, then $\text{LSB}(x) \in \text{RankC}(s)$

The candidates for the variable x while applying the above pattern are all the bit-vector variables that occur in the transition s . The query in the above pattern can be discharged using an SMT solver that provides support for bit-vector reasoning, and, in

particular, the LSB operator. (If the SMT solver does not provide first-class support for the LSB operator, then one can encode the LSB operator using bit-level manipulation as described in [31].) In the following we give example transitions whose bound can be computed using the above rule.

- $\text{RankC}(x' = x < 1 \wedge x \neq 0) = \{\text{LSB}(x)\}$
- $\text{RankC}(x' = x \& (x - 1) \wedge x \neq 0) = \{\text{LSB}(x)\}$

6.4 Data-structure Iteration Patterns

Iteration over data-structures or collections is quite common, and one standard way to iterate over a data-structure is to follow field dereferences until some designated object is reached. Such an iteration pattern can be captured by the following rule/lemma, where the function $\text{Dist}(x, z, f)$ denotes the number of field dereferences along field f required to reach z from x .

If $s \Rightarrow (x \neq z \wedge (\text{Dist}(x', z, f) < \text{Dist}(x, z, f)))$,
then $\text{Dist}(x, z, f) \in \text{RankC}(s)$.

The candidates for variables x, z and field f , while applying the above pattern are all variables \vec{x} and field names that occur in s . The query in the above pattern can be discharged using an SMT solver that implements a decision procedure for the theory of reachability and can reason about its cardinalities (e.g., [18]). Note, that $\text{Dist}(x, z, f)$ denotes the cardinality of the set of all nodes that are reachable from x before reaching z along field f . In the following we give example transitions whose ranking functions can be computed using an application of this pattern.

- $\text{RankC}(x \neq \text{Null} \wedge x' = x.\text{next}) = \{\text{Dist}(x, \text{Null}, \text{next})\}$
- $\text{RankC}(\text{Mem}' = \text{Update}(\text{Mem}, x.\text{next}, x.\text{next}.\text{next}) \wedge x \neq \text{Null} \wedge x.\text{next} \neq \text{Null}) = \{\text{Dist}(x, \text{Null}, \text{next})\}$

7. Bound Computation for Transition Systems

In this section, we show how to compute a bound for a transition system T .

If a transition system consists of a single transition s , then a bound for the transition system can be obtained simply from any ranking function r of the transition s using the following theorem.

THEOREM 14. *Let $r \in \text{Rank}(s)$. Then,*

$$\text{Bound}(s) = \text{Max}(0, r)$$

where the Max operator returns the maximum of its arguments.

PROOF: If the transition s is ever taken, then r denotes an upper bound on number of iterations of s (since, by our definition of a ranking function, transition s implies that r is bounded below by 0 and decreases by at least 1 in each iteration). The other case is when s is never executed (i.e., the number of iterations of s is 0). Combining these two cases, we obtain the result. \square

The significance of sanitizing the bound by applying the Max operator in Theorem 14 is illustrated in Example 20.

Obtaining a bound for a transition system consisting of multiple transitions is not as straight-forward. We cannot simply add the ranking functions of all individual transitions to obtain the bound for the transition system, since the interleaving of those transitions with each other can invalidate the decreasing measure of the ranking function. An alternative can be to define the notion of lexicographic ranking functions [6] or disjunctively well-founded ranking functions [28] for transition systems consisting of multiple transitions. Such an approach may sometimes work for proving termination, but would usually not be precise for yielding bounds.

For the purpose of precise bound computation, we distinguish between the different ways in which two transitions of a transition system can interact with each other. These cases (described in Sections 7.1, 7.2, and 7.3) allow for composing the ranking functions of the two transitions using one of three operators \max , sum , and product . These cases can be efficiently identified asking queries to SMT solvers.

7.1 Max Composition of Ranking Functions

The bound for a transition system consisting of two transitions $s_1 \vee s_2$ can be obtained by applying the Max operator to ranking functions for the individual transitions under cases when the transitions are either disjoint, or they decrease each other's ranking functions. In fact, the criterion is slightly more general, and is formalized in Theorem 16, which makes use of the following definition.

DEFINITION 15 (Cooperative-interference). *We say there is cooperative interference between transitions s_1 and s_2 through their ranking functions r_1 and r_2 , if any of the conditions below hold:*

- (Non-enabling condition) $s_1 \circ s_2 = \text{false}$.
- (Rank-decrease condition) $s_1 \Rightarrow r_2[x'/\vec{x}] \leq \text{Max}(r_1, r_2) - 1$.

We denote such a cooperative-interference by $\text{CI}(s_1, r_1, s_2, r_2)$.

THEOREM 16 (Proof Rule for Max-Composition).

Let $r_1 \in \text{RankC}(s_1)$ and $r_2 \in \text{RankC}(s_2)$. If $\text{CI}(s_1, r_1, s_2, r_2) \wedge \text{CI}(s_2, r_2, s_1, r_1)$, then

$$\text{Bound}(s_1 \vee s_2) = \text{Max}(0, r_1, r_2)$$

PROOF: We consider four cases below. (1) If both transitions s_1 and s_2 satisfy the non-enabling condition, then either only transition s_1 can execute or only transition s_2 can execute. Hence, the result. (2) If both transitions satisfy the rank-decrease condition, then it can be shown that $\text{Max}(r_1, r_2)$ is a ranking function for both the transitions s_1 and s_2 . Hence, the result. (3) Now suppose transition s_1 satisfies the non-enabling condition, while transition s_2 satisfies the rank-decrease condition. The only possibility is that a sequence of transitions s_2 is followed by a sequence of transitions s_1 . The result now follows from the fact that $\text{Max}(r_1, r_2)$ is a ranking function for s_2 , while r_1 is a ranking function for s_1 . (4) The last case is similar to (3). \square

EXAMPLE 17. Consider the transition system $s_1 \vee s_2$ from Figure 1(i) with the following 2 transitions:

$$\begin{aligned} s_1 &\stackrel{\text{def}}{=} (n' = n - 1 \wedge j < n \wedge j' \geq j \wedge i' = i) \\ s_2 &\stackrel{\text{def}}{=} (n' = n - 1 \wedge i < n - 2 \wedge i' \geq i + 1 \wedge j' \geq i + 2) \end{aligned}$$

We can compute $\text{RankC}(s_1) = \{n - j\}$ and $\text{RankC}(s_2) = \{n - i - 2\}$. We can prove $\text{CI}(s_1, n - j, s_2, n - i - 2)$ and $\text{CI}(s_1, n - i - 2, s_2, n - j)$. An application of the max-composition theorem yields a bound of $\text{Max}(0, n - i - 2, n - j)$ for the transition system $s_1 \vee s_2$.

7.2 Additive Composition of Ranking Functions

The bound for a transition system consisting of two transitions $s_1 \vee s_2$ can be obtained by adding together the ranking functions for the two transitions under cases when the transitions do not interfere with each other's ranking functions. To state this formally (Theorem 19), we first define the notion of *non-interference* of a transition with respect to the ranking function of another transition.

DEFINITION 18 (Non-interference). *We say that a transition s_1 does not interfere with the ranking function r_2 of another transition s_2 , if any of the following conditions hold:*

- (Non-enabling condition) $s_1 \circ s_2 = \text{false}$
- (Rank-preserving condition) $s_1 \Rightarrow (r_2[x'/\vec{x}] \leq r_2)$

We denote such a non-interference by $\text{NI}(s_1, s_2, r_2)$.

The following theorem holds. We use the notation $\text{Iter}(s)$ to denote the total number of iterations of transition s inside its transition system.

THEOREM 19 (Proof Rule for Additive-Composition). *Let $r_1 \in \text{RankC}(s_1)$, $r_2 \in \text{RankC}(s_2)$. If $\text{NI}(s_1, s_2, r_2) \wedge \text{NI}(s_2, s_1, r_1)$, then*

$$\begin{aligned} \text{Bound}(s_1 \vee s_2) &= \text{Iter}(s_1) + \text{Iter}(s_2), \text{ where} \\ \text{Iter}(s_1) &= \text{Max}(0, r_1) \\ \text{Iter}(s_2) &= \text{Max}(0, r_2) \end{aligned}$$

PROOF: The non-interference conditions $\text{NI}(s_2, s_1, r_1)$ ensure that the value of the ranking function r_1 for transition s_1 is not increased by any interleaving of transition s_2 . Hence, the total number of iterations of the transition s_1 is given by $\text{Max}(0, r_1)$ (based on an argument similar to that in proof of Theorem 14). Similarly, the total number of iterations of the transition s_2 is given by $\text{Max}(0, r_2)$. Hence, the result. \square

EXAMPLE 20. Consider the transition system $s_1 \vee s_2$ (obtained from the loop in the example program `Ex6` in Fig. 3) with the following 2 transitions:

$$\begin{aligned} s_1 &\stackrel{\text{def}}{=} z > x \wedge x < n \wedge x' = x + 1 \wedge \text{Same}(\{z, n\}) \\ s_2 &\stackrel{\text{def}}{=} z \leq x \wedge x < n \wedge z' = z + 1 \wedge \text{Same}(\{x, n\}) \end{aligned}$$

We can compute $\text{RankC}(s_1) = \{n - x\}$ and $\text{RankC}(s_2) = \{n - z\}$. We can prove $\text{NI}(s_1, s_2, n - z)$ and $\text{NI}(s_2, s_1, n - x)$. An application of additive-composition theorem yields a bound of $\text{Max}(0, n - x) + \text{Max}(0, n - z)$ for the transition system $s_1 \vee s_2$.

We now explain the importance of using the Max operators in the statement of Theorem 14 and Theorem 19. If we defined $\text{Iter}(s)$ to simply r instead of $\text{Max}(0, r)$, then we would incorrectly conclude the bound on the transition system $s_1 \vee s_2$ to be $2n - x - z$. This is incorrect because, for example, suppose that the transition system was executed in the initial state $n = 100, x = 0, z = 200$, then the expression $n - x - z$ evaluates to 0, while the transition system $s_1 \vee s_2$ executes for 100 iterations.

This example is also a good illustration of how our technique differs significantly from (and, in fact, provides a simpler alternative to) recently proposed techniques for proving termination [8] and loop bound analysis [16]. The control-flow refinement technique used in [16] unravels the exact interleaving pattern between the two transitions to conclude that s_1 and s_2 interleave in lock-steps, only after which it is able to derive the bound. In contrast, our proof rule stated in Theorem 19 only requires to establish the non-interference property between the two transitions. The principle of disjunctively well-founded ranking functions used in [8] requires computing the transitive closure of the transition system only to conclude a quadratic bound. In contrast, our proof rule stated in Theorem 19 does not require computing any transitive-closure, and is even able to obtain a precise linear bound. (The transitive-closure is required in our technique only to summarize any inner nested loops, which, however, are not present in the loop in the example program `Ex6`).

Observe, that the additive-composition and max-composition Theorems provide quite orthogonal proof-rules. The bound for the transition system in Example 17 can be computed using the max-composition Theorem, but not using the additive-composition Theorem. Similarly, the bound for the transition system in Example 20

can be computed using the additive-composition Theorem, but not using the max-composition Theorem.

7.3 Multiplicative Composition of Ranking Functions

If we cannot establish mutual cooperative-interference or mutual non-interference properties of two transitions, then it is still possible to compute bounds provided one of the transition satisfies the non-interference property. The bound in such a case is obtained by multiplying together the ranking functions for the two transitions, as made precise in the following theorem. This is a common case for bounding iterations of an inner loop when its iterators are re-initialized inside the outer loop leading to a multiplicative bound.

THEOREM 21 (Proof Rule for Multiplicative-Composition). *Let $r_1 \in \text{RankC}(s_1)$, and $r_2 \in \text{RankC}(s_2)$. If $\text{NI}(s_2, s_1, r_1)$, then*

$$\begin{aligned} \text{Bound}(s_1 \vee s_2) &= \text{Iter}(s_1) + \text{Iter}(s_2), \text{ where} \\ \text{Iter}(s_1) &= \text{Max}(0, r_1) \\ \text{Iter}(s_2) &= \text{Max}(0, r_2) + \text{Max}(0, u_2) \times \text{factor} \\ &\text{ where } \text{factor} = \text{Max}(0, r_1) \end{aligned}$$

where $u_2(\vec{x})$ denotes an upper bound on expression $r_2(\vec{x})$ in terms of \vec{x} as implied by $\text{TC}(s_1)$. For the special case when $(r_1 > 0) \wedge s_2$ is unsatisfiable, we can choose factor to be 1.

PROOF: From the non-interference condition $\text{NI}(s_2, s_1, r_1)$, we can conclude that $\text{Iter}(s_1) \leq \text{Max}(0, r_1)$ (the same argument as in the proof of Theorem 19). However, the same thing cannot be s_2 . Instead we observe that the maximum number of iterations of s_2 in between any two interleavings of s_1 is bounded above by $\text{Max}(0, u_2)$ (since the starting value of the ranking function r_2 is reset to u_2 by any execution of s_1). However, the number of iterations of s_2 before any interleaving of s_1 is still bounded by $\text{Max}(0, r_2)$. Hence, the total number of iterations of s_2 is bounded by $\text{Max}(0, r_2) + \text{Max}(0, u_2) \times \text{Max}(0, r_1)$. The special case follows from the observation that even though s_1 interferes with the ranking function r_2 of s_2 , it can interfere at most once since s_2 is enabled only after completion of all iterations (as opposed to somewhere in the middle) of s_1 . In other words, the worst-case possibility is a sequence of transitions s_2 , followed by a sequence of transitions s_1 , followed by a sequence of transitions s_2 . \square

EXAMPLE 22. Consider the transition system with the following two transitions s_1 and s_2 .

$$\begin{aligned} s_1 &\stackrel{\text{def}}{=} i' = i - 1 \wedge i > 0 \wedge j' = j - 1 \wedge j > 0 \wedge \text{Same}(\{k', m'\}) \\ s_2 &\stackrel{\text{def}}{=} j' = m \wedge k' = k - 1 \wedge k > 0 \wedge \text{Same}(\{i', m'\}) \end{aligned}$$

We can compute $\text{RankC}(s_1) = \{i, j\}$ and $\text{RankC}(s_2) = \{k\}$. We can prove $\text{NI}(s_1, s_2, k)$ and $\text{NI}(s_2, s_1, i)$. An application of additive-composition theorem yields a bound of $\text{Max}(0, i) + \text{Max}(0, k)$ for the transition system $s_1 \vee s_2$. An application of multiplicative-composition theorem yields an incomparable bound of $\text{Max}(0, j) + \text{Max}(0, m) \times \text{Max}(0, k)$.

7.4 Combining the Composition Rules

In this section, we discuss how to compute bounds for a transition system with multiple (including more than 2 transitions) by putting together the proof rules mentioned in Theorem 16, 19, and 21.

First, observe that an optimal way of applying the proof rules in additive-composition Theorem and multiplicative-composition Theorems (Theorem 19 and Theorem 21) is to compute the total number of iterations for each transition individually, and then sum them up together. The algorithm described in Figure 7 implements

```

ComputeBound( $\bigvee_{i=1}^n s_i$ )
1  for  $i \in \{1, \dots, n\}$ :  $\text{Iter}[s_i] := \perp$ ;
2  do {
3    for  $i \in \{1, \dots, n\}$  and  $r \in \text{RankC}(s_i)$ :
4       $J := \{j \mid \neg \text{NI}(s_j, s_i, r)\}$ ;
5      if  $(\text{Iter}[s_i] = \perp) \wedge (\forall j \in J : \text{Iter}[s_j] \neq \perp)$ 
6        factor := 0;
7        foreach  $j \in J$ : factor := factor +  $\text{Iter}[s_j]$ ;
8        Let  $u(\vec{x})$  be an upper bound on  $r[\vec{x}]/\vec{x}$ 
           as implied by  $\text{TC}(\bigvee_{j \neq i} s_j)$ .
9         $\text{Iter}[s_i] := \text{Max}(0, r) + \text{Max}(0, u) \times \text{factor}'$ ;
10 } while any change in Iter array;
11 if  $(\forall j \in \{1, \dots, n\} : \text{Iter}[s_j] \neq \perp)$ , return  $\sum_j \text{Iter}[s_j]$ ;
12 else return 'Potentially Unbounded'

```

Figure 7. Bound Computation for a Transition System $\bigvee_{i=1}^n s_i$ from ranking functions $\text{RankC}(s_i)$ of individual transitions.

such a strategy based on a simple extension of Theorem 19 and Theorem 21 to the case when a transition system contains more than 2 transitions. The algorithm iteratively computes an array Iter such that $\text{Iter}[s_i]$ denotes a bound on the total number of iterations of the transition s_i during any execution of the transition system $s_1 \vee \dots \vee s_n$. The array J at Line 4 contains the indices of all transitions that interfere with the ranking function r of transition s_i . If a bound on the total number of iterations of all those transitions is known (test on Line 5), then the iterations of s_i is obtained using a generalization of Theorems 19 and 21 (Line 9). A bound on the entire transition system is obtained by simply summing up the bound on the total number of iterations of the individual transitions (Line 11). For simplicity, we have presented the algorithm to output only one bound, but the algorithm can be easily extended to output multiple bounds by relaxing the condition $\text{Iter}[s_i] = \perp$ in Line 5 and by associating a set of bounds (as opposed to a single bound) with $\text{Iter}[s_i]$.

EXAMPLE 23. Consider the transition system $s_1 \vee s_2 \vee s_3$ (obtained from the loop in Ex7 in Figure 3) with the following 3 transitions:

$$\begin{aligned} s_1 &= j < n \wedge j < m \wedge j' = j + 1 \wedge 0 < n < m \wedge \text{Same}(\{n, m\}) \\ s_2 &= j > n \wedge j < m \wedge j' = j + 1 \wedge 0 < n < m \wedge \text{Same}(\{n, m\}) \\ s_3 &= j \geq m \wedge j' = 0 \wedge 0 < n < m \wedge \text{Same}(\{n, m\}) \end{aligned}$$

We can compute $\text{RankC}(s_1) = \{n - j, m - j\}$, $\text{RankC}(s_2) = \{m - j\}$, $\text{RankC}(s_3) = \{\text{Bool2Int}(j \geq m)\}$. Since $\text{NI}(s_1, s_2, m - j)$ and $\text{NI}(s_3, s_2, m - j)$, the algorithm in Figure 7 first computes $\text{Iter}[s_2] = \text{Max}(0, m - j)$. Using $\text{NI}(s_1, s_3, \text{Bool2Int}(j \geq m))$, the algorithm now computes $\text{Iter}[s_3] = \text{Bool2Int}(j \geq m) \times (1 + 1) \leq 2$. From $\text{NI}(s_2, s_1, n - j)$, the algorithm now computes $\text{Iter}[s_1] = \text{Max}(0, n - j) + \text{Max}(0, n) \times 2$. The algorithm now returns a total bound of $\text{Max}(0, m - j) + 2 + \text{Max}(0, n - j) + \text{Max}(0, n) \times 2$. This bound can be translated to a bound in terms of the inputs in the example program Ex7 by substituting $n + 1$ for j (as obtained from the initial state before the loop) to yield $m + 1 + n$, which is a factor of 2 away from the real bound of $m + 1$ (since $n < m$).

This example also illustrates how our technique differs significantly from (and, in fact, provides a simpler alternative to) recently proposed techniques for termination and loop bound analysis. The control-flow refinement technique used in [16] uses a sophisticated machinery to unravel the exact interleaving pattern between the three transitions (in particular, s_2 follows s_3 which in turn follows

s_1) and is able to obtain the exact bound of $m + 1$. In contrast, our proof rules yield a bound of $m + 1 + n$, but using a much simpler formalism. We do not know of any other technique (including [8, 19]) that can even prove termination of this example.

We now briefly discuss an extension to the above-described algorithm that also takes advantage of the proof rule in the max-composition Theorem (Theorem 16). Before running the algorithm, we extend $\text{RankC}(s)$ for any transition s with $\text{Max}(r, r')$, where $r \in \text{RankC}(s)$ and $r' \in \text{RankC}(s')$ for some other transition s' , provided $\text{Rank}(s, \text{Max}(r, r'))$ holds. This allows for an application of additive-composition Theorem to obtain an additive bound that is a constant factor of 2 away from what would have been obtainable from application of the max-composition Theorem (but much better than a multiplicative-bound). A more complete scheme that directly takes advantage of the max-composition Theorem is a bit involved and is left out for lack of space.

8. Experiments

We have implemented our proposed solution to the reachability-bound problem in C# using the Phoenix Compiler Infrastructure [26] and the SMT solver Z3 [1]. This implementation is part of a tool that computes symbolic complexity for procedures in .Net binaries. Below we present two different sets of experimental results that measure the effectiveness of various aspects of our solution.

8.1 Loop Bound Computation

We considered the problem of computing symbolic bounds on the number of loop iterations, which is an instance of the reachability-bound problem where the control location under consideration is the loop header. We chose mscorlib.dll (a .Net base-class library), which had 2185 loops, as our benchmark. Our tool analyzes these 2185 loops in less than 5 minutes and is able to compute bounds for 1677 loops. The problem of loop bound computation is especially challenging under the following two cases for which earlier techniques for bound computation do not perform as well.

Case 1: Iterations of outer loops depending on inner loops (examples of the kind described in Figure 2). There were 113 such loops out of the total 2185 loops. The key idea of our paper to address such challenges is to replace the inner loops by their transitive-closure that preserves required relationships between the inputs and outputs of the loop. The *effectiveness of our transitive closure computation algorithm* is illustrated by the fact that our success ratio for such cases (80 out of 113, i.e., 70%) is similar to our overall success ratio (1677 out of 2185, i.e., 76%).

Case 2: Loop bound computation for nested loops. The challenge here is to compute precise amortized bounds on the total number of iterations of those loops, as opposed to the number of iterations per iteration of the immediate outer loop (the latter is an easier problem than the former). This is the same issue as exemplified by the example in Figure 1. There were 250 such loops out of the total 2185 loops. Unfortunately, we cannot evaluate the precision of our bounds automatically. As described in Section 3, the problem of computing a precision-witness for a given symbolic bound is an orthogonal problem that we are currently working on. Instead, we manually investigated the generated bounds for most of these loops and found all these bounds to be precise (according to Definition 2). This points out the *effectiveness of our bound-computation algorithm* based on the three proof rules presented in Section 7.

Another interesting statistic is the distribution of the number of transitions generated for each loop, as shown in Figure 8. The small number of transitions validates the design choice behind our *transition system generation algorithm* that enumerates all paths

# Transitions	1	2	3	4	5	6	7	8	9	≥ 10
# Loops	1561	224	107	44	25	11	9	5	8	191

Figure 8. Number of loops for respective number of transitions.

between two program points (in order not to loose any precision) after slicing has been performed.

Out of the 508 loops for which we failed to compute a bound, the failure for 503 loops is attributed to not being able to compute ranking functions for some transition in the transition system corresponding to the loop. There were two main causes. (i) Our implementation is intra-procedural, meaning that our transition system generation algorithm fails when the value of loop iterators gets modified because of procedure calls. This problem can be addressed by simply inlining the procedure, provided there are no recursive calls. (ii) Of the various proof rules described in Section 6, we only implemented those corresponding to arithmetic and boolean iteration patterns, while several transitions were iterating using field dereferences or bit-vector manipulation. A sound handling of field dereferences would require use of an alias analysis. A more optimistic way to read this statistic is to observe the *effectiveness of the proof-rule based technique for finding ranking functions*: a handful of patterns are sufficient to compute ranking functions for transitions arising in 76% of the examples.

There were only 5 cases (out of 1682 cases) for which we were able to compute a ranking function for each transition, but were not able to compute a bound for the transition system. This points out the *effectiveness of our proof rules for bound computation* from composition of ranking functions of individual transitions.

8.2 Disjunctive Invariant Computation

We also evaluated the effectiveness of our transitive closure algorithm on a variety of benchmark examples chosen by recent state-of-the-art papers on computing disjunctive invariants. Figure 9 describes these four examples that have been used as flagship examples to motivate new techniques for proving non-trivial safety assertions. Proving validity of the assertions in all these examples requires disjunctive loop invariants. It turns out that the required disjunctive invariant for each of these examples satisfies the convexity-like assumption, and hence can be discovered by our transitive closure algorithm in Figure 6. We adapt our algorithm slightly to take advantage of the initial condition (as is done by all the other approaches) by initializing s'_1 to $\text{Init} \wedge \text{Id}$ at Line 2, instead of only Id since Init is known at the beginning of each loop. This allows our algorithm to establish the desired assertion using a disjunctive invariant with fewer disjuncts. (For a more detailed discussion on this adaptation, see the end of this section).

Given that the number of disjuncts in the desired transitive closure is 2 for all examples, and that the number of transitions in the transition system represented by the loop is either 2 or 3, the total number of possibilities for the map σ is 16 or 64 respectively. Hence, by trying out all possible maps, the algorithm in Figure 6 can discover the desired disjunctive invariants.

Instead, we experimented with a heuristic for dynamic construction of map σ that we found to be effective for all examples. We choose $m = 1$ and initialize s'_1 to $\text{Init} \wedge \text{Id}$. We maintain a partial map σ that is completely undefined to start with, and use the following heuristic to construct σ on the fly. For each choice of (i, j) on Line 4 in the algorithm, if $\sigma(j, i)$ is undefined, we compute $s = s'_j \circ s_i$ in the abstract domain. If s is not equal to false, then we use a semantic-merging criterion to find any k such that s is close to an existing disjunct s'_k and define $\sigma(j, i)$ to be k . If no such k exists, we increase m by 1 and define $\sigma(j, i)$ to the new value of m . The semantic-merging criterion that we used for our experiments was one that checks agreements on variable equalities (as opposed to the more general inequality relationships expressible

Original Example	Various Details
Gopan and Reps 06. P. 3, F. 1 $x:=0, y:=0;$ while (*) if ($x \leq 50$) $y++;$ else $y--;$ if ($y<0$) break; $x++;$ assert($x=102$)	$(x \leq 50 \wedge y+1 \geq 0 \wedge y' = y+1 \wedge x' = x+1)_{s_1}$ $\vee (x > 50 \wedge y-1 \geq 0 \wedge y' = y-1 \wedge x' = x+1)_{s_2}$ Init $\equiv x = 0 \wedge y = 0$ $(0 \leq x' \leq 51 \wedge x' = y')_{s'_1}$ $\vee (52 \leq x' \leq 102 \wedge x' + y' = 102)_{s'_2}$ $\delta = 1, \sigma = \{(1, 1) \mapsto 1, (1, 2) \mapsto 2, (2, 2) \mapsto 2, (2, 1) \mapsto 1\}$
Beyer et al. 07. P. 306, F. 4. $x:=0; y:=50;$ while ($x<100$) if ($x<50$) $x++;$ else $x++; y++;$ assert($y=100$);	$(x \leq 50 \wedge x' = x+1 \wedge y' = y)_{s_1}$ $\vee (51 \leq x \leq 100 \wedge x' = x+1 \wedge y' = y+1)_{s_2}$ Init $\equiv x = 0 \wedge y = 50$ $(0 \leq x' \leq 50 \wedge y' = 50)_{s'_1}$ $\vee (51 \leq x' \leq 100 \wedge x' = y')_{s'_2}$ $\delta = 1, \sigma = \{(1, 1) \mapsto 1, (1, 2) \mapsto 2, (2, 2) \mapsto 2, (2, 1) \mapsto 1\}$
Gulavani et al. 06. P. 5, F. 3. Henzinger et al. 02. P. 2, F. 1. lock:=0; assume($x \neq y$) while ($x \neq y$) lock := 1; $x := y;$ if (*) lock := 0; $y++;$ assert(lock = 1);	$(x \neq y \wedge lock' = 1 \wedge x' = y \wedge y' = y)_{s_1}$ $\vee (x \neq y \wedge lock' = 0 \wedge x' = y \wedge y' = y+1)_{s_2}$ Init $\equiv x \neq y \wedge lock = 0$ $(x' = y' \wedge lock' = 1)_{s'_1}$ $\vee (x' + 1 = y' \wedge lock' = 0)_{s'_2}$ $\delta = 1, \sigma = \{(1, 1) \mapsto 1, (2, 1) \mapsto 1, (1, 2) \mapsto 2, (2, 2) \mapsto 2\}$
Popeea and Chin 06. P. 2 $x := 0; upd := 0;$ while ($x < N$) if (*) $l := x; upd := 1;$ $x++;$ assert($upd = 1 \Rightarrow 0 \leq l < N$);	$(x < N \wedge x' = x+1 \wedge l' = l \wedge upd' = upd)_{s_1}$ $\vee (x < N \wedge x' = x+1 \wedge l' = x \wedge upd' = 1)_{s_2}$ Init $\equiv x = 0 \wedge upd = 0$ $(x' \geq 0 \wedge l' = l \wedge upd' = 0 \wedge N' = N)_{s'_1}$ $\vee (x' \geq 1 \wedge upd' = 1 \wedge N' = N \wedge 0 \leq l' < N)_{s'_2}$ $\delta = 1, \sigma = \{(1, 1) \mapsto 1, (1, 2) \mapsto 2, (2, 2) \mapsto 2, (2, 1) \mapsto 2\}$

Figure 9. Prominent disjunctive invariant challenges from recent literature. Entries in 2nd column show the following details in that order: transition-system representation of the loop, initial condition Init, transitive closure of the transition-system required to prove the assertion, and the convexity-witness (δ, σ) .

in the octagon domain [24] used by our prototype implementation). This heuristic is an excellent example of combining the strengths of semantic-merging criterions in the light of the importance of having a static syntactic merging criterion as suggested by Theorem 12 (if we do not want to iterate over all maps σ). We implemented this heuristic and our prototype implementation is able to validate the assertion in each of the examples in less than 0.2sec.

We now return to the discussion on what would happen if we do not adapt our algorithm to make use of the initial condition Init while computing a loop summary. We can still prove the desired assertion, but the required transitive closure would consist of more disjuncts, and would involve elements from a numerical domain richer than the octagon abstract domain. For example, for the first example, we would require the following disjunctive invariant:

$$\begin{aligned}
& (Id)_{s'_1} \vee (x \leq 50 \wedge x' \leq 51 \wedge x' - x = y' - y)_{s'_2} \\
& \vee (x \geq 51 \wedge x' \geq 52 \wedge x' - x = y' - y)_{s'_3} \\
& \vee (x \leq 50 \wedge x' \geq 52 \wedge 102 - x' - x = y - y')_{s'_4}
\end{aligned}$$

Observe, that the above invariant again satisfies the convexity-like assumption, where a convexity-witness σ is as follows: $\sigma = \{(1, 1) \mapsto 2, (2, 1) \mapsto 2, (3, 1) \mapsto 3, (4, 1) \mapsto 4, (1, 2) \mapsto 3, (2, 2) \mapsto 4, (3, 2) \mapsto 3, (4, 2) \mapsto 4\}$. Hence, our approach can be used to discover this invariant. In contrast, none of the techniques presented for the respective examples can analyze the loops in such a *modular* setting where the initial condition is not

initially known. Further discussion on the use of our technique for modular analysis is beyond the scope of this paper.

9. Comparison with Related Work

Disjunctive Invariant Generation A variety of techniques exist to lift classical abstract domains (like intervals, octagons [24], and polyhedra [10]), which typically infer conjunctive invariants, to the powerset extension or some approximation of it for discovering disjunctive invariants [20, 29, 13, 14]. These techniques address the hardness inherent in this problem by proposing various semantic-merging heuristics. In contrast, we present a result that calls for working with a static syntactic merge criterion under the convexity-like assumption (which appears to be satisfied by the benchmark examples).

Some syntactic techniques based on program restriction [5] or control-flow refinement [16] have also been suggested for discovering disjunctive invariants. These can be viewed as instantiations of our more general framework based on a convexity-witness σ .

Symbolic Bound Generation There is recent work on generating symbolic bounds on the number of loop iterations [16, 19, 15, 2], but none of these techniques directly addresses the more general problem of reachability-bound that we introduce in our paper. Our solution reduces the reachability-bound problem to the problem of computing bounds of an outer loop, but one whose iterations are influenced by inner loops. None of [16, 19, 15, 2] directly address the challenge of computing bounds for such loops, and hence would fail to compute bounds for most of the examples presented in the paper. In contrast, our technique can compute bounds for all the motivating examples presented in [16, 19, 15, 2].

[19] would fail to compute bounds for the example programs Ex1, Ex3, Ex4, Ex5, Ex7 because the invariants required for establishing bounds on the counters are disjunctive. (It can only compute bounds for Ex2 and Ex6.) The multiplicative counter instrumentation strategies that are meant to alleviate the problem of computing disjunctive invariants do not help in this case because there is only one back-edge for the outer loop and only one counter can be instrumented.

[16] would fail to compute bounds for Ex1, Ex3, Ex4, Ex5 for the same reason of requiring disjunctive invariants for performing the desired reasoning on inner loops. (It can only compute bounds for Ex2, Ex6 and Ex7.) The control-flow refinement strategy is meant to alleviate the problem of computing disjunctive invariants, but it does not help in any of these cases since the control-flow is already refined, and it cannot be refined any further.

[15] requires user annotations to identify interesting non-linear and disjunctive expressions to compute bounds for transition systems with multiple transitions. We address these challenges by means of novel proof rules. However, the technique described in [15] can be used in a synergistic manner with our technique, in particular, as an extension to the pattern-matching based technique to compute bounds/ranking-functions for single transitions.

[2] computes bounds by generating recurrence relations and then deriving a closed form expression for the maximum size of the unfoldings of the recurrence relations into trees. Since they do not precisely summarize inner loops, they cannot handle loops where the inner loop changes the iterators of the outer loop as in the example programs Ex2, Ex3, Ex4, Ex5. Also, they can't handle examples Ex6, Ex7 and are unable to compute the amortized complexity as in the example program Ex1.

We report the first implementation of symbolic bound generation for .Net binaries, while [19, 16, 15] and [2] implemented bound generation for C++ and Java programs respectively. Hence, we only provide analytical (not experimental) comparison with these techniques. Quite significantly, our implementation scales to large pro-

grams, while [19, 15, 2] have been applied to only small benchmarks.

[12] computes symbolic bounds by curve-fitting timing data obtained from profiling. Their technique has the advantage of measuring real time in seconds for a representative workload, but does not provide worst-case bounds. There is a large body of work on estimating worst case execution time (WCET) in the embedded and real-time systems community [30, 32]. WCET research is largely orthogonal, focused on distinguishing between the complexity of different code-paths and low-level modeling of architectural features such as caches, branch prediction, instruction pipelines. For establishing loop bounds, WCET techniques either require user annotation, or use simple techniques based on pattern matching or simple numerical analysis. These WCET techniques cannot compute bounds for most of the examples considered in this paper.

[11] presents a type system for the certification of resource bounds (once they are provided by the programmer). In contrast, we *infer* bounds. [22] uses linear programming to infer bounds for functional programs, but they are restricted to computing only linear bounds.

Termination Analysis There has been a large body of work on proving termination of programs and the standard approach used has been that of finding ranking functions. We also use ranking functions to compute bounds, but our focus is on finding *precise* ranking functions, using composition by Max or $+$ operators if possible, that can yield precise symbolic bounds. Bounds can also be obtained from the standard lexicographic ranking functions or disjunctively well-founded ranking relations [8], but only using multiplicative-composition, which is imprecise compared to the bounds that can be obtained from max- or additive- composition.

In fact, our proof rules can also be regarded as an alternative new technique for proving termination. For example, the recently proposed approach based on variance assertions or disjunctively well-founded ranking relations cannot be used to prove termination of the loop in the example program Ex7, while our technique can.

There is superficial similarity between termination techniques based on computing variance assertions [4], transition invariants [28] and disjunctively well-founded ranking relations [8] in that they also summarize relationships between two different visits to a control location, and often require disjunctive invariants. However, there are two key technical differences: (a) Our technique requires computing relationships between two immediate visits to a control location, while the approach based on transition invariants or variance assertions requires computing relationships between *any* two visits to a control location. (b) Our technique requires use of disjunctive invariants only to summarize nested loops.

10. Future Work and Conclusion

This paper defined and motivated the reachability-bound problem. The paper also presented a solution to the reachability-bound problem in the context of non-recursive and sequential programs. The next technical challenge is to address the reachability-bound problem in context of recursive procedures and concurrent execution.

On the applications side, we are working on integrating the proposed solution to the reachability-bound problem with other specific techniques to provide an integrated solution for resource bound analysis in some contexts such as memory bound analysis, and active-task graph size analysis in asynchronous programs.

References

- [1] Z3 Theorem Prover. research.microsoft.com/projects/Z3/.
- [2] E. Albert, P. Arenas, S. Genaim, and G. Puebla. Automatic inference of upper bounds for recurrence relations in cost analysis. In *SAS*, 2008.
- [3] A. M. Ben-Amram. Size-change termination, monotonicity constraints and ranking functions. In *CAV*, pages 109–123, 2009.
- [4] J. Berdine, A. Chawdhary, B. Cook, D. Distefano, and P. O’Hearn. Variance analyses from invariance analyses. In *POPL*, 2007.
- [5] D. Beyer, T. A. Henzinger, R. Majumdar, and A. Rybalchenko. Path invariants. In *PLDI*, pages 300–309, 2007.
- [6] A. Bradley, Z. Manna, and H. Sipma. Termination of polynomial programs. In *VMCAI*, 2005.
- [7] S. Chaudhuri, S. Gulwani, and R. Lubliner. Continuity analysis of programs. In *POPL*, 2010.
- [8] B. Cook, A. Podelski, and A. Rybalchenko. Termination proofs for systems code. In *PLDI*, pages 415–426, 2006.
- [9] P. Cousot and R. Cousot. Abstract Interpretation: A Unified Lattice Model for Static Analysis of Programs by Construction or Approximation of Fixpoints. In *POPL*, pages 238–252, 1977.
- [10] P. Cousot and N. Halbwachs. Automatic Discovery of Linear Constraints among Variables of a Program. In *POPL*, 1978.
- [11] K. Crary and S. Weirich. Resource bound certification. In *POPL ’00*.
- [12] S. Goldsmith, A. Aiken, and D. S. Wilkerson. Measuring empirical computational complexity. In *ESEC/SIGSOFT FSE*, 2007.
- [13] D. Gopan and T. W. Reps. Lookahead widening. In *CAV*, 2006.
- [14] D. Gopan and T. W. Reps. Guided static analysis. In *SAS*, 2007.
- [15] B. S. Gulwani and S. Gulwani. A numerical abstract domain based on expression abstraction and max operator with application in timing analysis. In *CAV*, pages 370–384, 2008.
- [16] S. Gulwani, S. Jain, and E. Koskinen. Control-flow refinement and progress invariants for bound analysis. In *PLDI*, 2009.
- [17] S. Gulwani and S. Juvekar. Bound analysis using backward symbolic execution. Technical report, Oct 2009.
- [18] S. Gulwani, T. Lev-Ami, and M. Sagiv. A combination framework for tracking partition sizes. In *POPL*, 2009.
- [19] S. Gulwani, K. Mehra, and T. Chilimbi. Speed: precise and efficient static estimation of program computational complexity. In *POPL ’09*.
- [20] M. Handjieva and S. Tzolovski. Refining static analyses by trace-based partitioning using control flow. In *SAS*, pages 200–214, 1998.
- [21] T. Henzinger. From boolean to quantitative system specifications, keynote. In *1st Workshop on Quantitative Analysis of Software*. <http://research.microsoft.com/users/sumitg/qa09/keynote.pdf>, 2009.
- [22] S. Jost, H. Loidl, K. Hammond, and M. Hofmann. Static determination of quant. resource usage for higher-order programs. In *POPL ’10*.
- [23] P. Malacaria. Assessing security threats of looping constructs. In *POPL*, pages 225–235, 2007.
- [24] A. Miné. The octagon abstract domain. In *WCRE*, 2001.
- [25] S. S. Muchnick. *Advanced Compiler Design and Implementation*. Morgan Kaufmann, 1997.
- [26] Microsoft Phoenix Compiler, research.microsoft.com/phoenix/.
- [27] A. Podelski and A. Rybalchenko. A complete method for the synthesis of linear ranking functions. In *VMCAI’04*.
- [28] A. Podelski and A. Rybalchenko. Transition invariants. In *LICS ’04*.
- [29] C. Popeea and W.-N. Chin. Inferring disjunctive postconditions. In *ASIAN*, pages 331–345, 2006.
- [30] A. Prantl, J. Knoop, M. Schordan, and M. Triska. Constraint solving for high-level wcet analysis. *CoRR*, 2009.
- [31] H. S. Warren. *Hacker’s Delight*. Addison-Wesley Longman Publishing Co., Inc., Boston, MA, USA, 2002.
- [32] R. Wilhelm, J. Engblom, A. Ermedahl, N. Holsti, S. Thesing, D. Whalley, G. Bernat, C. Ferdinand, R. Heckmann, F. Mueller, I. Puaut, P. Puschner, J. Staschulat, and P. Stenström. The Determination of Worst-Case Execution Times—Overview of the Methods and Survey of Tools. In *ACM Transactions on Embedded Computing Systems (TECS)*, 2007.